

1 Design a metrics reporting system

1.1 Question

Talk me through how you would design a system to create and report metrics for user actions and other events occurring on an Android app.

Assume the metrics will be sent to a SQL database through a REST API which you control. [The backend design is out of scope.](#)

Examples of metrics your system should support include:

- counts (e.g. button clicks)
- timers (e.g. time taken for a network call)
- aggregate statistics (e.g. data usage in a given time period)

1.2 Expected clarifying questions

- What attributes should be in each metric? What kind of queries are expected to be performed?
 - Should be able to query according to deviceId, platform version, device model
 - Should be able to query according to which screen/workflow/state the app was in when the metric was emitted
- Do the metrics need to be uploaded in real-time? What kind of delay is acceptable?
- Does the app support offline usage? If so, is there a limit to the batch size?
- Is there a minimum acceptable network strength before the metric can be uploaded?

1.3 Design considerations

- Storage/persistence
- Threading. Events can be generated in main thread but need to get uploaded in the background
- Network connectivity, backend unavailability - need a way to store the metrics
- Network/data usage
- Battery usage

1.4 Criteria

	Raises the bar (SDE II)	Meets the bar (SDE II)	Lowens the bar (SDE II)
Storage/persistence	<ul style="list-style-type: none">• understanding the tradeoffs and coming up with multiple options	<ul style="list-style-type: none">• Design involves persisting metrics as local File, AtomicFile or local DB	<ul style="list-style-type: none">• Does not account for offline or poor

	<ul style="list-style-type: none"> Picking one approach over the other. and explaining how each approach would work. 	<p>asynchronously when the metric is generated</p> <ul style="list-style-type: none"> Candidate accounts for staleness (e.g. metric persisted locally for 1 week without a successful upload ...) 	<p>connectivity cases, which necessitate storing the metric locally until they can be successfully uploaded</p> <ul style="list-style-type: none"> Proposes to store the metrics either as instance state Bundle (which does not consider complete process death) or in Shared Pref (which does not support durable disk writes)
Upload	<ul style="list-style-type: none"> Without prompting, design supports cases when fill rate is occasionally higher than drain rate (i.e. metrics being generated faster than they can be uploaded) 	<ul style="list-style-type: none"> With some prompting, design supports occasions when fill rate is occasionally higher than drain rate by setting a max batch size 	<ul style="list-style-type: none"> Attempts to correct imbalance solely by increasing frequency of uploads. Disregards issues with battery usage and network

			performance
Concurrency	<ul style="list-style-type: none"> Leverages DB transactions to handle concurrency, without any custom locking/synchronization For multiple tiers, supports a solution that uploads metrics for each tier independently of one another 	<ul style="list-style-type: none"> Without prompting, recognizes that metrics can be generated in UI thread but uploaded in background thread Handles concurrency through custom locking/synchronization blocks using existing ios libraries instead of creating your own 	<ul style="list-style-type: none"> Needs to be prompted to account for threading Proposes to use AsyncTask (Android specific) Design would cause main thread to freeze (App Not Responding , i.e. ANR) Created a queue and gave it a high priority.
Network usage	<ul style="list-style-type: none"> Candidate offers industry-standard solutions to minimize data usage per upload (e.g. GZIP compression, protobufs) Distinction why they are using Protobufs vs JSON Supports retry in the case of server failure, and also handles 	<ul style="list-style-type: none"> Attempts to reduce data usage by uploading metrics as a batch Supports a retry mechanism in the case of server failures With some prompting, differentiates between various failure responses. e.g. considers how to 	<ul style="list-style-type: none"> Does not support batching (i.e. design involves uploading each metric one by one) Does not support retry in the case of

	<p>scenarios with multiple consistent server failures</p> <ul style="list-style-type: none"> • With minimal/no prompting, differentiates between various failure responses. e.g. considers how to treat HTTP 400 (BAD REQUEST) vs. HTTP 500 (SERVER UNAVAILABLE). What is the customer impact? • How do they use the bad request? • back off algorithm. 	<p>treat HTTP 400 (BAD REQUEST) vs. HTTP 500 (SERVER UNAVAILABLE)</p>	<p>server failures</p>
Battery usage	<ul style="list-style-type: none"> • Candidate offers solutions that employ Android framework/OS support for minimizing battery usage (e.g. WorkManager, JobScheduler) • Adaptiveness-plugged in vs not plugged in • Wifi versus cellular 	<ul style="list-style-type: none"> • Attempts to reduce battery usage by uploading metrics periodically at suitably large intervals (e.g. greater than 1min) 	<ul style="list-style-type: none"> • Upload as soon as metric is created
Configurability	<ul style="list-style-type: none"> • Without prompting, candidate's solution supports multiple tiers of metrics (e.g. CRITICAL vs. NORMAL vs. DEBUG), with each tier's storage and upload policies 	<ul style="list-style-type: none"> • Supports multiple tiers with some prompting by splitting the metrics up into multiple queues or tables. 	<ul style="list-style-type: none"> • Naive support for multiple tiers (e.g. iterate through queue to find and

	<p>independently configurable according to network strength, available storage etc.</p> <ul style="list-style-type: none"> Without prompting, candidate offers solutions to make the system configurable for multiple apps setting up feedback loops 	<ul style="list-style-type: none"> Turn on and turn off features 	<p>upload CRITICAL metrics and upload, then iterate through same queue for NORMAL metrics, etc.)</p>
General design	<ul style="list-style-type: none"> When deciding how to account for cases where app crashes just before/after uploading metrics, candidate favors solution that minimizes 'lost' metrics Design supports multiple Upload policies that provide unified handling of fill-rate issue, battery usage, network usage Modularity : making library Extensible fields 	<ul style="list-style-type: none"> Asking clarifying questions for metrics using the available framework Adding basic fields 	<ul style="list-style-type: none"> Confuses metrics reporting with logging, even after repeated prompts Proposes to use an existing 3rd-party solution and is unable to explain its internals